

Springer,
Nested Class Modularity in Squeak/Smalltalk

Nested Class Modularity in Squeak/Smalltalk

by

Matthias Springer

A thesis submitted to the
Hasso-Plattner-Institute
at the University of Potsdam, Germany
in partial fulfillment of the requirements for the degree of
MASTER OF SCIENCE IN IT-SYSTEMS ENGINEERING

Supervisors

Prof. Dr. Robert Hirschfeld

Software Architecture Group
Hasso-Plattner-Institute
University of Potsdam, Germany

July 17, 2015

Abstract

The english abstract.

Zusammenfassung

Die Zusammenfassung auf deutsch.

Acknowledgments

I owe everything to my cat.

Contents

1. Introduction	1
1.1. Modularity	1
1.2. The Squeak Programming Language	1
1.3. Outline of this Work	1
2. Modularity Problems in Squeak	3
2.1. Class Name Clashes	3
2.2. Dependency Managment	3
2.3. Readability and Understandability	3
2.4. Code Reuse	3
3. Related Work	5
3.1. Class Name Clashes	5
3.1.1. Namespaces/Packages	5
3.1.2. Squeak Environments	5
3.1.3. Newspeak Modules	5
3.2. Dependency Management	5
3.2.1. Java Class Loader	5
3.2.2. Separate Compilation	5
3.2.3. External Configuration in Newspeak	5
3.3. Readability and Understandability	5
3.3.1. Smalltalk Packages	5
3.3.2. Hierarchical Decomposition	5
3.3.3. Information Hiding with Interfaces	5
3.4. Code Reuse	5
3.4.1. Multiple Inheritance	5
3.4.2. Mixins	5
3.4.3. Traits	5
4. Nested Class Modularity in Squeak	7
4.1. Nested Classes	7
4.2. Accessing the Lexical Scope	8
4.3. Parameterized Classes	11
5. Implementation	15
5.1. Meta Model and Instantiation	15
5.2. Anonymous Classes and Subclass Generation	15
5.3. thisOuter and thisScope	15

Contents

5.4. Class Updates	15
5.5. Integration in Squeak	15
5.5.1. Module Repository	15
5.5.2. IDE Support	15
5.5.3. Debugger	15
6. Use Cases	17
6.1. Avoiding Clasds Name Clashes	17
6.2. Module Versioning	17
6.3. Dependency Management	17
6.4. Readability and Understandability	17
6.5. Mixin Modularity with Parameterized Classes	17
6.6. Class Generator Pattern	17
6.7. Extension Methods	17
7. Future Work	19
7.1. Class as Instance-side Members	19
7.2. Bytecode Transformation instead of Recompilation	19
7.3. Squeak Integration	19
A. First Unimportant stuff.	23

List of Figures

4.1. Nested Classes Example	8
4.2. Keywords for access to superclass and lexical scope.	9
4.3. Binding of enclosing to method's lexical scope.	10
4.4. Implementation of Mixins with Nested Classes	12

List of Tables

List of Listings

List of Abbreviations

API application programming interface

1. Introduction

1.1. Modularity

1.2. The Squeak Programming Language

1.3. Outline of this Work

2. Modularity Problems in Squeak

2.1. Class Name Clashes

flat global namespace

2.2. Dependency Managment

two application require the same module but in different versions

2.3. Readability and Understandability

only one level of grouping: package

2.4. Code Reuse

share behavior among multiple classes

3. Related Work

3.1. Class Name Clashes

3.1.1. Namespaces/Packages

VisualWorks, Java, Ruby, Python

3.1.2. Squeak Environments

3.1.3. Newspeak Modules

3.2. Dependency Management

3.2.1. Java Class Loader

3.2.2. Separate Compilation

3.2.3. External Configuration in Newspeak

3.3. Readability and Understandability

3.3.1. Smalltalk Packages

3.3.2. Hierarchical Decomposition

Java, Python, Ruby, Newspeak, ...

3.3.3. Information Hiding with Interfaces

3.4. Code Reuse

3.4.1. Multiple Inheritance

3.4.2. Mixins

Ruby Modules, Python Multiple Inheritance, Newspeak, Jigsaw

3.4.3. Traits

Squeak implementation

4. Nested Class Modularity in Squeak

In this chapter, we describe the main concept of our work: classes as class members. Similar concepts are part of programming languages like Java, Ruby, Python, and Newspeak. Our concept follows closely the Newspeak notion of nested classes, but without making invasive changes to the Smalltalk programming language.

4.1. Nested Classes

In Smalltalk, every object is an instance of a class, defining the object's instance variables and the messages it understands. Consequently, a class is also an instance of its so-called meta class. Every meta class is an instance of Metaclass. In the remainder of this work, we denote the meta class of a class *C* by *C class*. Every Smalltalk image has a `globals` dictionary, mapping symbols to class objects, so that references to classes can be resolved at compile time. This implies that all references to classes are early bound.

Our system extends the Smalltalk class organization as follows: in addition to regular methods, we introduce the concept of *class generator methods*. Such a method generates a class and is associated with a set *I* of instance methods and a set *C* of class methods. Whenever the method is invoked, the system first executes the method body, then adds *I* to resulting class and *C* to resulting meta class, and finally returns the resulting class. For performance reasons, our system also caches the result, meaning that a class is not generated twice.

Details Class generator methods are only allowed as class-side methods. Instance-side class generator methods seem to provide neglectable benefits and make the implementation of our system more complicated. We provide an in-depth explanation of instance-side class generator methods in the Section 7.1.

A class generated by a class generator method is anonymous: it is not listed in the `globals` dictionary and can only be referenced using message sends to its enclosing class¹. Consequently, its name is a concatenation of all class names on the path from the top-level class to class in question.

Notation and Example Figure 4.1 shows an example of nested classes in Squeak. *A* is a top-level class, i.e., it is part of the `globals` dictionary and known everywhere in the system; it can be referenced by just writing the identifier *A*. *A* has one instance

¹It can also be referenced by sending the `cClass` method to one of its instances

4. Nested Class Modularity in Squeak

method `m2` and two class methods `m1` and `B`. In accordance with UML notation, class-side method selectors are underlined.

A `class»B` is a class generator method that is associated with a set of instance methods `{}` and a set of class methods `{foo, B}`. The name of the class it generates is `A B`, which is in that case also a valid Smalltalk code expression that evaluates to the generated class. A `class»B class»C` is a class generator method that generates `A B C`. Note, that we use the `»` notation to not only reference methods but also the classes they generate, in case they are class generator methods.




Figure 4.1.: Nested Classes Example

4.2. Accessing the Lexical Scope

It is sometimes necessary to access a method's lexical scope (i.e., the enclosing classes), in order to send messages to enclosing classes. For this reason, our system introduces new keywords, in addition to `self` and `super`, which are already present in every Smalltalk dialect. Figure 4.2 gives an overview of all method lookup-related keywords in the system.

self Keyword This keyword is used make a message send within an object. The receiver is the same object as the sender and the lookup starts at the (polymorphic) class of the receiver. If that class does not provide a corresponding method, the lookup continues in the superclass hierarchy. If no class in the superclass has a corresponding method, a `MethodNotUnderstood` error is raised.

super Keyword This keyword is also used to make a message send within an object. Again, the receiver is the same object as the sender, but the lookup starts at

4.2. Accessing the Lexical Scope




Figure 4.2.: Keywords for access to superclass and lexical scope.

superclass of the sender's method class. Note, that `super` is bound to the superclass of the method class, not the superclass of the receiver's class.

enclosing Keyword This keyword is used to make a message send to the class that contains the current current. Consider, for example, that we want to send a message `foo` to class `A B` within `A class»B class»C»m5` in Figure 4.1. Either one of the following two statements works in this case.

- `A B foo.`
- `enclosing foo.`

`enclosing` is a keyword that evaluates to the method owner's enclosing class upon method compilation. Note, that `enclosing` is bound to the method's lexical scope, not the receiver's lexical scope.

Figure 4.3 illustrates how `enclosing` is bound. In `B1 class»C class»bar1`, `enclosing` is bound to `B1`. In contrast, `B2 class»C class»bar2` binds `enclosing` to `B2`. Consequently, `B1 C bar1` calls `B1 foo` and so does `B2 C bar1`, even though the receiver of `bar1` is an instance of `B2 C class` and not `B1 C class` in the latter case. Note, that `B2 C bar2` calls `B2 foo`, because `bar2`'s lexically enclosing class is `B2`.

Note, that `enclosing` can be used for meta programming purposes; however, it should be avoided in general. Our system also provides a `scope` keyword that should be used instead.

enclosing Method In addition to `enclosing`, every class in the system has a method `enclosing` that returns the enclosing class of the receiver², making it possible to send messages to enclosing classes which are more than one level away.

²The enclosing class of an object that is not a class is its class' enclosing class.

4. Nested Class Modularity in Squeak




Figure 4.3.: Binding of enclosing to method's lexical scope.

If, for example, in Figure 4.1, A `class»B class»C»bar` wants to send the message `m1` to A, either one of the following two statements works.

- A `m1`.
- `enclosing enclosing m1`.

Again, the method enclosing should be avoided in general, but is useful to implement parts of our system with code written in the system itself.

outer Keyword The method enclosing can be used to traverse the the lexical scope of a class. Arbitrarily many enclosing sends can be chained, as long as the respective receiver still has an enclosing class and is, therefore, not a top-level class. Arguably, this can result in verbose and complicated code, and is at the very least questionable with regards to the law of demeter.

In addition to enclosing, the system provides the outer keyword, bound to the method's lexical scope. Whenever a message is sent to outer, the message is first interpreted as a send to enclosing. If that message send fails, the message is sent to enclosing enclosing, and, eventually, to the top-level class, if no other class in the lexical scope understands the message. If even that message send is not understood, the selector is looked up in the `globals` dictionary. If the selector is absent, a `MessageNotUnderstood` error is raised.

outer is similar to super, with the difference that outer does a horizontal lookup (lexical scope) and super does a vertical lookup (superclass chain). Note, that messages sent to outer are sent to an object different from `self`.

scope Keyword This keyword combines super and outer: a message sent to scope is first treated as a `self` send. If the message is not understood, it is treated as an outer send.

Our system essentially first looks up the methods in `self`, then in the superclass hierarchy, and then in the lexical scope. This is also how the method lookup in

4.3. Parameterized Classes

erence:
W.R.: New-
ot: Pro-
on the
p. 109 –
ingerVer-
9), in
oe-Based
nming:
ts, Lan-
and Ap-
ns, Noble,
ari and
editors

Java works, also known as *comb semantics*. Newspeak uses a different lookup: it first looks for a method in the receiver's class, then in the lexical scope, and finally in superclass hierarchy [1].

The statement enclosing enclosing m1 in the previous example can also be written as scope m1. If the method m1 would now be moved to its enclosing class (if it had one), the lookup would still succeed. However, scope exposes the risk of accidentally capturing method names in superclasses or the lexical chain.

Implicit scope Receiver In our system, references to globals are in fact message sends with scope as implicit receiver. This makes it easier for Smalltalk programmers to write code in our system, even if they do not know about enclosing and scope. It also makes the code less verbose and easier to read.

Whenever code references an identifier that is not a temporary variable, not an instance variable, and not a *special* object ³, the compiler replaces that identifier with a message send to scope.

Consider, for example, that we want to reference class A B within A class»B class»C»bar in Figure 4.1. Either one of the following two statements works in this case.

- A B.
- enclosing.
- enclosing enclosing B.
- outer B.
- scope B.
- B.

In this example, we used the implicit scope receiver for class lookup, which is in our opinion the most useful case. However, any method in self, the lexical scope, or the superclass hierarchy can in fact be looked up this way. One can argue that this is bad practice and should be forbidden for methods that are not class generator methods. However, it is allowed in Newspeak and other programming languages like Java, and seems to work well, as long as the programmer is aware of how the method lookup works. Note, that only unary messages can have an implicit scope receiver, since we would have to change the Smalltalk syntax otherwise.

4.3. Parameterized Classes

All examples shown in the previous section use unparameterized classes, i.e., class generator methods are always unary. Class generator method can, however, also have binary selectors or selectors with a higher arity. For memory conservation reasons, these classes are then no longer cached.

³self, super, thisContext, scope, outer, enclosing

4. Nested Class Modularity in Squeak

Mixins Parameterized classes can be used to build mixins. Mixins are not a special feature of this system: they are an application of our system and come for free by just having class nesting as described in the previous sections; they are an immediate consequence of parameterized classes. A mixin is a function that takes as an input a class and outputs a subclass with additional behavior, i.e., it is a class transformer.

A mixin can be implemented by writing a class generator method with one parameter which is the input class. The method creates a subclass of that input class and returns it. Associated with that parameterized class generator methods is a set of instance-side methods and a set of class-side methods. These are the methods that will be added when applying the mixin.

Recursive Mixin Application A mixin can make sure another mixin is applied upon its application. This is done creating a subclass of a mixin application in the class generator method. Consequently, the system first creates a subclass of the base class, adds the methods of the inner mixin, then creates a subclass of the resulting class, and finally adds the methods of the outer mixin.

Example Figure 4.4 shows an example of parameterized classes and how they can be used to build mixins.

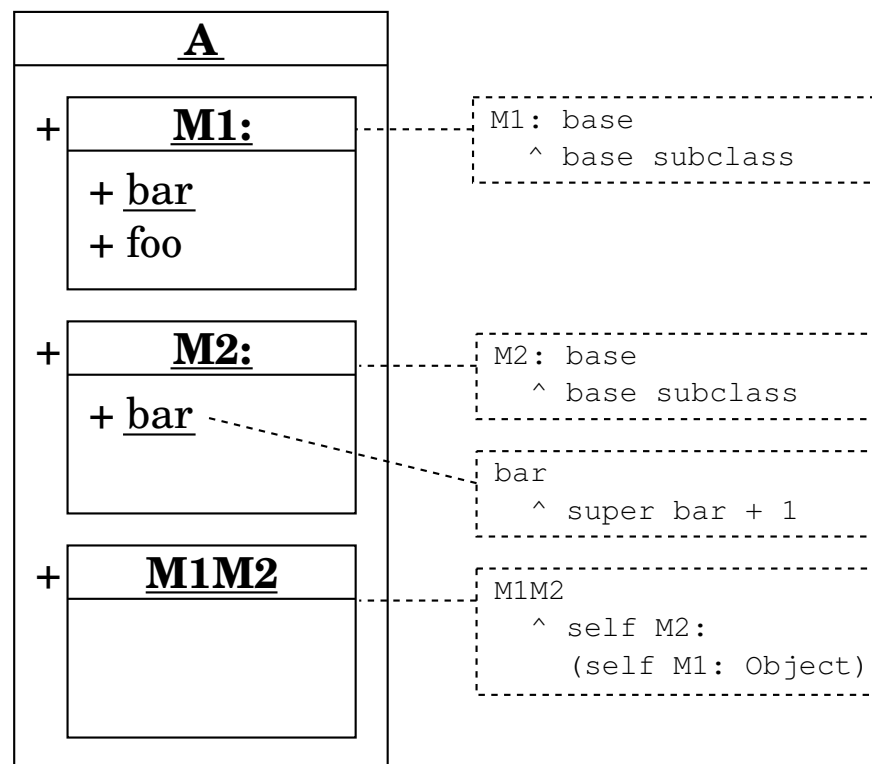


Figure 4.4.: Implementation of Mixins with Nested Classes

4.3. Parameterized Classes

Two class generator methods `A M1:` and `A M2:` are defined, which take as input a base class and output a subclass with additional behavior. `A M1M2` is an application of both both mixins. `A M1M2`'s superclass is *some* `A M2:`, whose superclass is *some* `A M1:`, whose superclass is `Object`. Note, that `A M1:` and `A M2:` are not specific classes: we use this notation as a name for *some* application of `A class»M1:` and `A class»M2:`, respectively. Therefore, even if two classes have the same name, they are not necessarily the same class if they names contain a colon.

Note, that evaluating `A M1: Object` multiple times returns different class object, since parameterized classes are not cached. However, `A M1M2` is cached, because it is a unary method. Therefore, calling `A M1M2` multiple times always returns the same class object.

The notation used in `A class»M1M2` can be a bit confusing at first. That method first applies `A M1:` to `Object`, and then `A M2:`; however, in the source code, `A M2:` appears before `A M1:`. For readability reasons, and to support more features like pre-include hooks and post-include hooks, we present the Class Generator Pattern in Section 6.6.

5. Implementation

5.1. Meta Model and Instantiation

accessor methods and generator methods, lazy initialization

5.2. Anonymous Classes and Subclass Generation

5.3. `thisOuter` and `thisScope`

5.4. Class Updates

using instances weak array on specification

5.5. Integration in Squeak

5.5.1. Module Repository

replacement for Smalltalk dict

5.5.2. IDE Support

works in workspace, test runner. How to write tests? New system browser

5.5.3. Debugger

shows slightly different code (`thisContext` automatically inserted, generator methods)

6. Use Cases

6.1. Avoiding Clasds Name Clashes

example: multiple games, all providing game class

6.2. Module Versioning

example: multiple versions of the same module in the same image

6.3. Dependency Management

example: app that requires a specific version (defining local alias), then all access goes through that method. example: external configuration of modules with parameterized classes (alternative to dependency injection)

6.4. Readability and Understandability

example: large project, where parts of the code can now be understood, given that we have hierarchical nesting. could be an example where grouping according to multiple criteria is needed (would result in $n \times m$ packages)

6.5. Mixin Modularity with Parameterized Classes

6.6. Class Generator Pattern

better syntax for class generators

6.7. Extension Methods

better way is needed (e.g., class boxes, refinements, COP). return already existing class in generator method

7. Future Work

7.1. Class as Instance-side Members

7.2. Bytecode Transformation instead of Recompilation

7.3. Squeak Integration

Bibliography

- [1] Gilad Bracha, Peter von der Ahé, Vassili Bykov, Yaron Kashi, William Maddox, and Eliot Miranda. “Modules as Objects in Newspeak”. English. In: *ECOOP 2010 – Object-Oriented Programming*. Ed. by Theo D’Hondt. Vol. 6183. Lecture Notes in Computer Science. Springer Berlin Heidelberg, 2010, pp. 405–428. ISBN: 978-3-642-14106-5. DOI: [10.1007/978-3-642-14107-2_20](https://doi.org/10.1007/978-3-642-14107-2_20).

Appendix A.

First Unimportant stuff.

Hello, here is some text without a meaning. This text should show what a printed text will look like at this place. If you read this text, you will get no information. Really? Is there no information? Is there a difference between this text and some nonsense like “Huardest gefburn”? Kjift – not at all! A blind text like this gives you information about the selected font, how the letters are written and an impression of the look. This text should contain all letters of the alphabet and it should be written in of the original language. There is no need for special content, but the length of words should match the language.

Eigenständigkeitserklärung

Hiermit versichere ich, dass ich die vorliegende Arbeit selbständig verfasst sowie keine anderen Quellen und Hilfsmittel als die angegebenen benutzt habe.

Potsdam, den 17. Juli 2015

Matthias Springer

Todo list

- Add reference: Smith, W.R.: NewtonScript: Prototypes on the Palm, pp. 109 – 139. SpringerVerlag (1999), in Prototype-Based Programming: Concepts, Languages and Applications, Noble, Taivalsaari and Moore, editors 11